



ASK TO DOC

Your friendly health buddy on Telegram
Get advice on symptoms, treatments, and wellness

*Made by:
Chiara Musso*

OVERVIEW



Training



Evaluation



Security



Telegram



TRAINING DETAILS

DATASET AND MODEL



RUSLANMV/AI-MEDICAL-CHATBOT DATASET

This is an experimental Dataset designed to run a Medical Chatbot It contains at least 250k dialogues between a Patient and a Doctor.




MICROSOFT/DIALOGPT-SMALL

DialoGPT is a SOTA large-scale pretrained dialogue response generation model for multiturn conversations

TRAINING SETUP

- **Dataset Preparation:** Split it into 90% training and 10% testing.
- **Tokenization:** Used AutoTokenizer from Hugging Face
- **Device Configuration:** Leveraged GPU for faster training.
- **Optimizer and Scheduler:** Implemented AdamW optimizer and linear learning rate scheduler.
- **Data Preparation:** Created a custom function to format patient-doctor conversations.
- **DataLoader:** Set up for training and evaluation with a batch size of 8.



TRAINING LOOP

- Used gradient accumulation to manage large batch sizes and reduce memory usage.
- Implemented mixed precision training with GradScaler for efficiency.
- Trained on random 10% subsets for quicker iterations.
- Tracked loss curves and metrics.





EVALUATION

METRICS

Why Evaluate?:

- To ensure the model's responses are accurate and relevant.
- To identify areas for improvement and validate the model's performance.



Perplexity

Measures how well the model predicts the next token in a sequence. Lower values indicate better performance.

BLUE score

Evaluates the similarity between generated responses and reference responses

ROGUE score

ROUGE-1: Measures overlap of unigrams between generated and reference texts.

ROUGE-2: Measures overlap of bigrams.

ROUGE-L: Measures the longest common subsequence overlap.

RESULTS

Perplexity	30.31
Average Loss	3.41
Average BLUE score	0.270
ROGUE 1	0.516
ROGUE 2	0.498
ROUGUE L	0.510



EXISTING MEDICAL CHATBOT

1. Ada Health

- Combines rules and AI
- High user satisfaction
- Accurate symptom triage
- No detailed metrics available

2. Babylon Health

- Uses deep learning and NLP
- Accuracy comparable to human doctors in some cases
- Mixed results from independent evaluations

3. Mayo Clinic Chatbot

- Uses extensive medical knowledge and rules
- Known for trustworthiness and accuracy
- Less advanced AI capabilities

Comparison Insights:

1. **Customization:** Fine-tuning for specific medical areas outperforms general models.
2. **Model Size:** Larger models capture more context and nuances.
3. **Data Augmentation:** Diverse, larger datasets improve response quality.
4. **Decoding Strategies:** Top-k sampling enhances response coherence and relevance.
5. **Comprehensive Evaluation:** Metrics like METEOR, CIDEr, and human feedback provide a better performance overview.



SECURITY

SECURITY EVALUATION

- Data Protection and Privacy
- Authentication and Authorization
- Secure Communication:
- Session Management
- Input Validation
- Storage and Handling of Personal Data
- Compliance with Regulations
- Threat Detection and Response
- API Security
- User and Developer Awareness



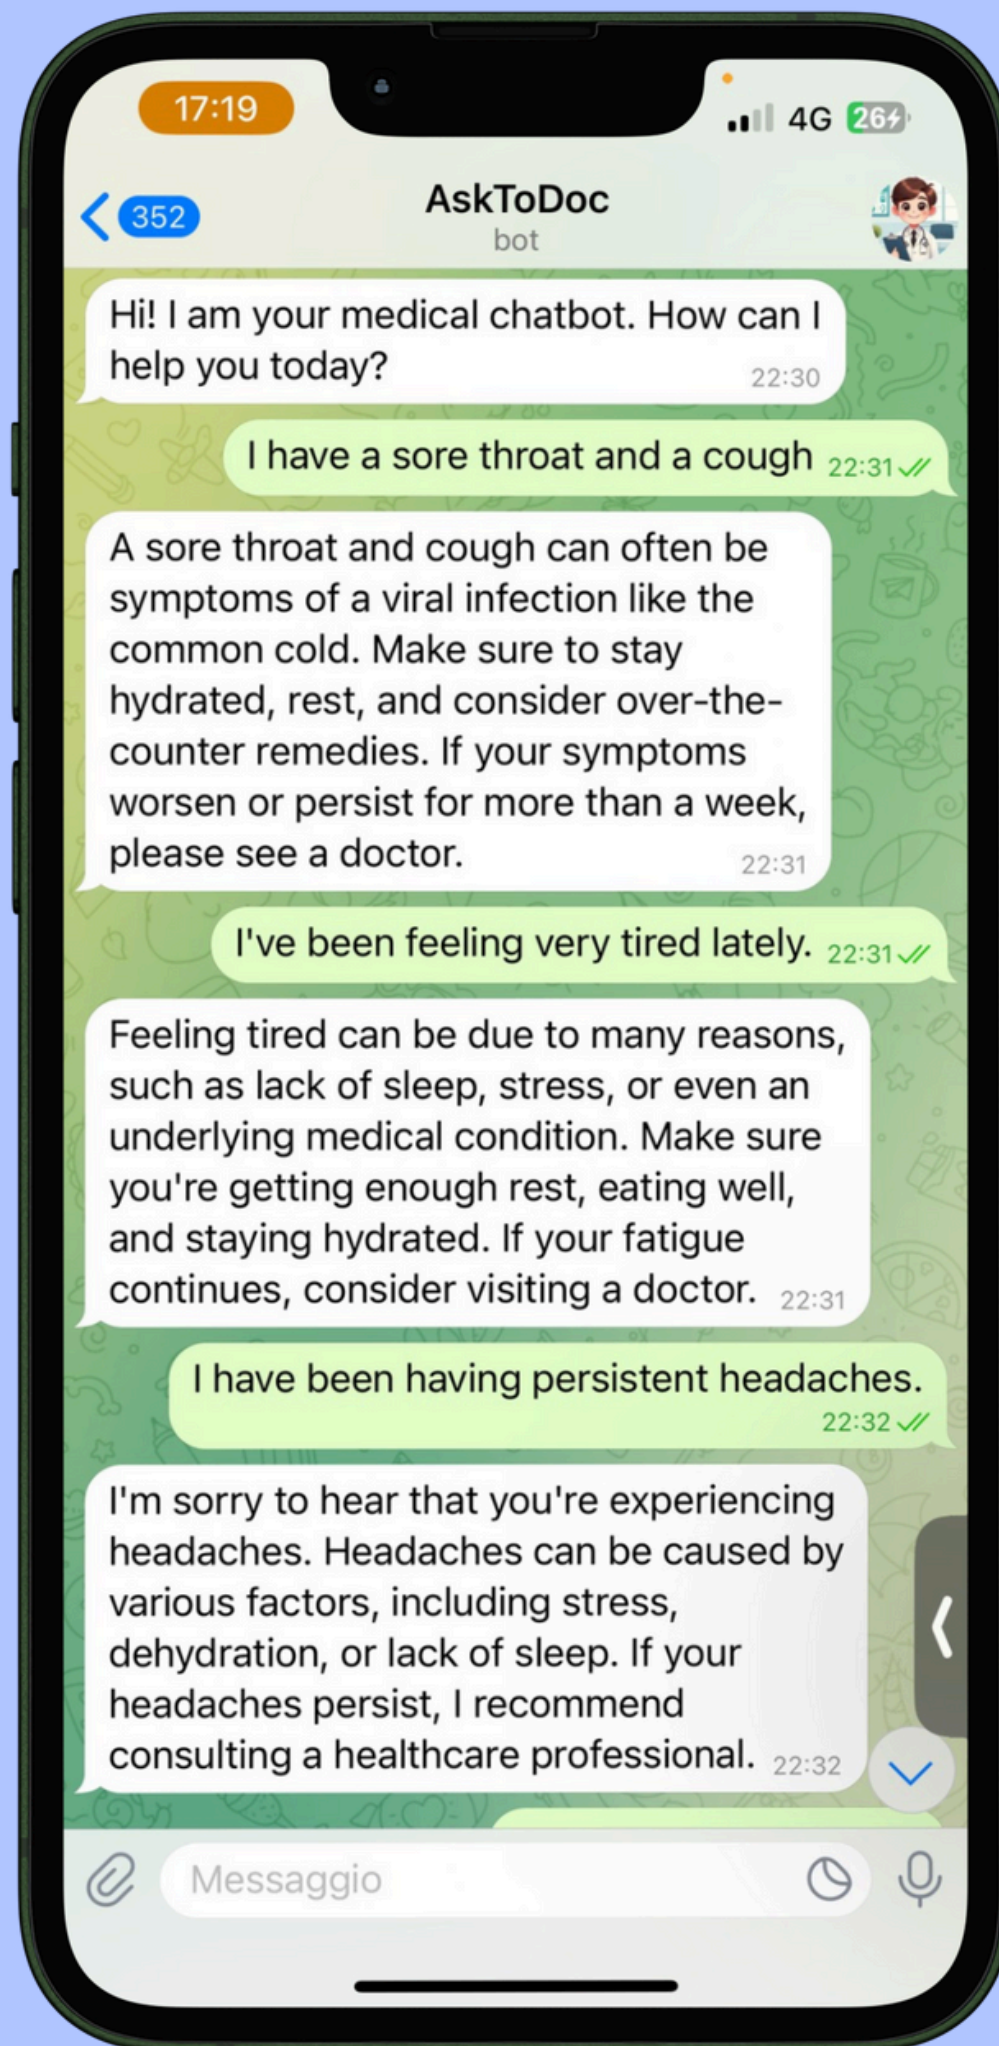
SPECIFIC TECHNIQUES AND PRACTICAL IMPLEMENTATIONS

- Regular Expressions and Pattern Matching: To detect and mask personal data.
- Audit Trails: Maintain detailed audit trails of user interactions and data access.
- Secure Backend Infrastructure: Ensure backend servers are secure and data is encrypted.
- Compliance with Data Retention Policies: Align data retention with legal requirements and user agreements.
- User Consent Management: Clearly inform users about data usage and obtain explicit consent.





TELEGRAM



@AskToDoc

HOW IT WORKS?

1. Bot Initialization

- The bot is set up with a token from the Telegram API.
- Command and message handlers are registered to manage different types of user interactions.

2. Handling Commands

- **Start Command:** When a user sends /start, the bot sends a greeting message.
- **Message Response:** The bot processes user messages, checks for keywords, and generates appropriate responses.

3. Generating Responses

- User input is tokenized and fed into the model.
- The model generates a response based on the input.



```
def generate_response(patient_text):
    # Tokenize the input text
    input_text = f"Patient: {patient_text} Doctor:"
    input_ids = tokenizer.encode(input_text, return_tensors='pt').to(device)

    # Generate a response from the model
    with torch.no_grad():
        output_ids = model.generate([
            input_ids,
            max_length=150,
            num_beams=5,
            no_repeat_ngram_size=3,
            early_stopping=True,
            pad_token_id=tokenizer.eos_token_id,
            temperature=0.7,
            top_k=50,
            top_p=0.95,
            do_sample=True # Enable sampling
        ])
```


FUTURE DEVELOPMENT



- Advanced AI Techniques
- Dataset Expansion
- User Feedback Integration
- Multilingual Support
- Enhanced Security and Privacy
- Integration with Additional Platforms
- Utilizing Advanced GPU Resources
- Continuous Performance Monitoring and Improvement
- Collaboration with Medical Professionals

**THANK YOU
FOR YOUR
ATTENTION**

